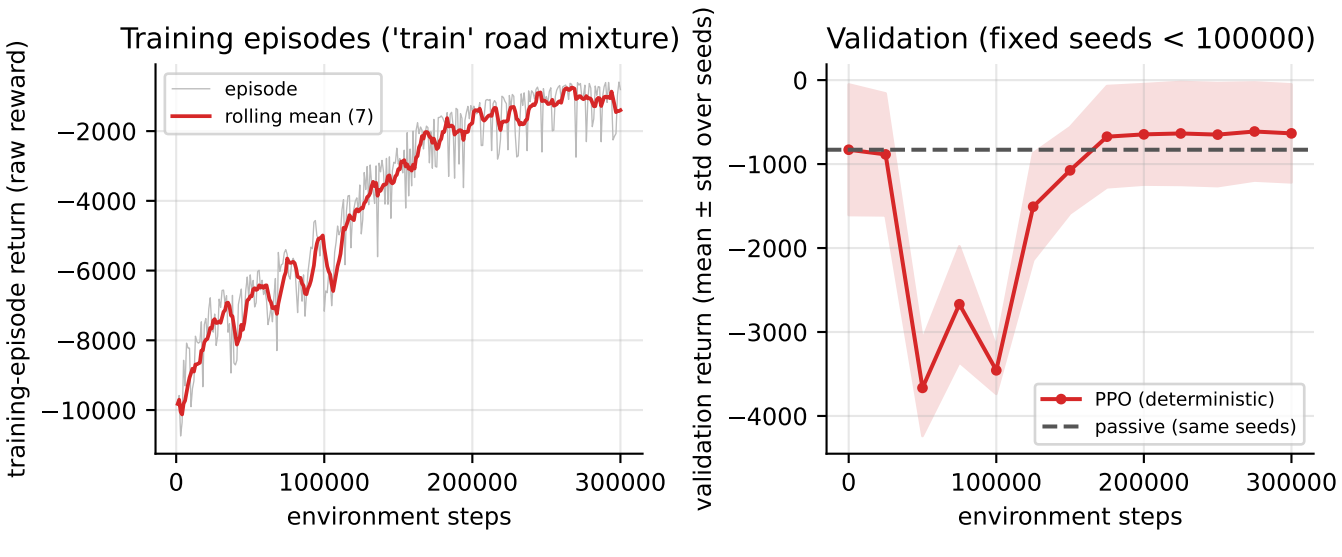
Training episodes ('train' road mixture)
episode
rolling mean (7)
training-episode return (raw reward)
−2000
−4000
−6000
−8000
−10000
0
100000
200000
300000
environment steps
Validation (fixed seeds < 100000)
validation return (mean ± std over seeds)
0
−1000
−2000
−3000
−4000
PPO (deterministic)
passive (same seeds)
0
100000
200000
300000
environment steps